




An Introduction to SPARQL

Cen Xi Toh

Knowledge Engineer


cenxi.toh@oxfordsemantic.tech



F1 Dashboard with Knowledge Graphs?



Controls and filters for the scoring system.



Statistics about the data used to form the results.

This matches the filters.

Try it for yourself!

<http://f1.rdfox.tech>

Requirements




A.
Get an RDFox License

[Free license here](#)



B.
Download RDFox (& unzip)

[Download here](#)



C.
Download the class materials from Github

[Download here](#)



D.
Download VS Code

[Download here](#)

Or your IDE of choice.
Setup steps won't be covered.

Objectives



The Foundations

- What is a database?
- What is a triple?
- What is a Knowledge Graph?
- What is Semantic Reasoning?
- What is RDFox?

Setting up RDFox

- License and executable
- VS Code
- REST endpoint and UI
- Create a data store
- Import data

Querying with SPARQL

- Basics of SPARQL
- Useful queries to explore
- Aggregates
- Negation
- Filters
- Binds
- Optionals

Please feel free to ask questions at any time! Exercises are scattered throughout. Links are provided for reference.

<https://docs.oxfordsemantic.tech/>

Objectives



The Foundations

- What is a database?
- What is a triple?
- What is a Knowledge Graph?
- What is Semantic Reasoning?
- What is RDFox?

Setting up RDFox

- License and executable
- VS Code
- REST endpoint and UI
- Create a data store
- Import data

Querying with SPARQL

- Basics of SPARQL
- Useful queries to explore
- Aggregates
- Negation
- Filters
- Binds
- Optionals

What is a database?




In practical terms, a database is a collection of information



We can retrieve this information on-demand with queries



RDFOx is a Knowledge Graph database



What is a Knowledge Graph?




A Knowledge Graph database stores information as nodes and relationships




Just like relational databases, we retrieve answers via queries



Unlike relational databases, there is no fixed schema



What is a triple?



A triple is a unit of data in a Knowledge Graph: one fact




Triples have 3 parts: 2 nodes and their relationship



Nodes can either be an entity or a value



Relationships (predicates) are directional—they point from the subject to the object



How to query a Knowledge Graph?




Knowledge Graphs retrieve data using SPARQL queries



SPARQL finds patterns in the data by describing the triples we're trying to match



In SPARQL, variables are declared with a question mark



What is Semantic Reasoning?



Semantic Reasoning (knowledge-based AI) enriches a Knowledge Graph by adding new information



This combines knowledge and data to answer more complex questions beyond queries alone




It drastically improves performance of queries by doing the hard work ahead of time



Reasoning updates automatically with new data

Rule:

Count the number of founders



What can Reasoning do for you?



```
1 #Query 11
2 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
3
4 # Drivers with their win percentage, ordered by win percentage
5 SELECT ?forename ?surname ?raceCount ?raceWins ?percentage
6 WHERE {
7   # First get the drivers
8   ?driver :driver_forename ?forename ;
9           :driver_surname ?surname .
10
11 # Then get the race count for each driver with an innery query...
12   {SELECT ?driver (COUNT(?race) AS ?raceCount)
13     WHERE {
14       ?result :result_driver ?driver ;
15              :result_race ?race .
16     }
17   GROUP BY ?driver}
18
19 # ... and get the *win* count for each driver with another inner query.
20   {SELECT ?driver (COUNT(?race) AS ?raceWins)
21     WHERE {
22       ?result :result_driver ?driver ;
23              :result_race ?race ;
24              :result_positionOrder 1 .
25     }
26   GROUP BY ?driver}
27
28 # Finally use the two aggregate variables to compute a percentage
29 # with the BIND keyword.
30 BIND(?raceWins/?raceCount AS ?percentage)
31
32 }
33 ORDER BY DESC(?percentage)
```

Fetches 108 answers in 0.011 s. ✓

**Datalog
rules**

```
1 #Query 19
2 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
3
4 SELECT ?forename ?surname ?percentage
5 WHERE {
6   ?driver :driver_forename ?forename;
7           :driver_surname ?surname;
8           :hasWinPercentage ?percentage;
9           :hasRaceCount ?count .
10 }
11 ORDER BY DESC(?percentage)
```

Fetches 108 answers in 0.003 s. ✓

What is RDFox[®]?



RDFox is a Knowledge Graph database
and Semantic Reasoning Engine



**Incremental
Reasoning**



**Speed at
scale**



**From device
to cloud**



**Oxford
University**

Objectives



The Foundations

- ✓ What is a database?
- ✓ What is a triple?
- ✓ What is a Knowledge Graph?
- ✓ What is Semantic Reasoning?
- ✓ What is RDFox?


Setting up RDFox

- License and executable
- VS Code
- REST endpoint and UI
- Create a data store
- Import data

Querying with SPARQL

- Basics of SPARQL
- Useful queries to explore
- Aggregates
- Negation
- Filters
- Binds
- Optionals

Organizing your files



1. Make sure to put your **RDFox folder INSIDE** the **RDFox Workshop folder**

Your workshop folder will be called **'RDFoxWorkshop'** if you cloned it (rather than downloading it)

2. Then drop your **RDFox license INSIDE** the **RDFox folder**

Without this setup you will need to use different file paths in the commands we provide.




Opening the terminal in VS Code



1 Open VS Code


2 Open the workshop folder with **File > Open Folder**

Select the folder directly containing the 'data', 'rules', 'queries' folders etc.




If the folder name appears twice, repeat step 2 and open the inner-most folder


3 Open a new terminal with **Terminal > New Terminal**



› E.g. MacOS (zsh)



› E.g. Windows (PowerShell)



This path is unique to you



Starting RDFox

- 1 In the terminal, from the 'RDFox Workshop' working directory, run the relevant command for the version of RDFox you downloaded:

You will need to edit the command if you are not using RDFox v7.6 or are running on LINUX

› MacOS ARM

```
./RDFox-macOS-arm64-7.6/RDFox sandbox
```

› MacOS INTEL

```
./RDFox-macOS-x86_64-7.6/RDFox sandbox
```


The Windows command depends on the shell used – shown right of the terminal window

› Windows (CMD)

```
RDFox-win64-x86_64-7.6\RDFox.exe sandbox
```

› Windows (PowerShell)

```
& RDFox-win64-x86_64-7.6\RDFox sandbox
```



- 2 The RDFox server should now be running!

If not, see next slide for common fixes

```
A new server connection was opened as role 'guest' and stored with name 'sc1'.
```


Common fixes



File path needs quotes and/or backslash

All systems

command not found *is not recognised* *permission denied* *is a directory*

Problem: Some terminals handle slashes and special characters differently

Fix: Surround the file path (excluding 'sandbox') with quotation marks "" and, on Windows, replace forward-slashes '/' with back-slashes '\'

Unzipping creates a nested folder

Predominantly Windows

The system cannot find the path specified. *The term is not recognized*

Problem: Unzipping a folder sometimes puts the contents in a folder of its own name

Fix: Use the inner 'RDFoxWorkshop' as the working directory (open VS code at the inner folder)

Check your folder structure

All systems

Ensure your folder structure follows RDFoxWorkshop > RDFox-macOS-arm64-7.6 > RDFox.lic

Check your working directory and spelling

All systems

Ensure your terminal is running from the RDFoxWorkshop folder and your commands are correct

Open the RDFox Console




- 1 In the terminal, run: `endpoint start`

```
> endpoint start
The REST endpoint was successfully started at port number/service name 12110 with 8 threads.
WARNING: The RDX endpoint is running with no transport layer security (TLS). This could allow attackers to steal
information including role passwords or authorization tokens. See the endpoint.channel variable and related
variables in the description of the RDX endpoint for details of how to set up TLS.
> 
```

- 2 The RDX endpoint is now running, so we can use the web UI

- 3 Open a browser and go to: `localhost:12110/console`





Loading data into RDFox – Create data store

1

Welcome to the RDFox console

This console allows you to connect with and use your RDFox instance. Connect with a data store to continue.

Select an existing data store

Or create a new data store

2

Create data store

A data store is a queryable container in RDFox that stores facts, rules, and OWL axioms.

Name data store

Advanced settings

Make sure to call your data store **f1** so that the exercise links work.

3

Data store created


Now we can import our data...

Learn more about getting started in the Console here:

<https://docs.oxfordsemantic.tech/getting-started.html#getting-started-with-the-web-console>



Loading data into RDFox – Importing files



Add content ×

Add facts, rules, or OWL axioms to your data store.

Type to select graph

Adding to the Default Graph


Drag files to upload

or

Supported formats: application/n-triples, text/turtle, application/trig, application/x.datalog, text/owl-functional.

Selected files

 upTo2020.ttl ×

☐ Update prefixes ☐ Off ☒ On

4 Remember to set
'Update prefixes' to on!
This imports the prefixes
defined in the file.



The data for today's class


Formula 1 data from the fan-run Ergast database

- Drivers
- Race results
- Races
- Constructors

Ergast Developer API

Query Results


Query Details				
Series	Page	Results		
f1	1 of 86	857		
Driver Table				
Driver Name	Permanent Number	Nationality	DOB	Information
Carlo Abate		Italian	1932-07-10	Biography
George Abecassis		British	1913-03-21	Biography
Kenny Acheson		British	1957-11-27	Biography



Exploring the data



[Click to open the console.](#)



The graph shows a small sample of the data.

Classes are green,
instances of those
classes are purple, and
properties of those
instances have no
boxes.

Objectives



The Foundations

- ✓ What is a database?
- ✓ What is a triple?
- ✓ What is a Knowledge Graph?
- ✓ What is Semantic Reasoning?
- ✓ What is RDFox?

Setting up RDFox

- ✓ License and executable
- ✓ VS Code
- ✓ REST endpoint and UI
- ✓ Create a data store
- ✓ Import data

Querying with SPARQL

- Basics of SPARQL
- Useful queries to explore
- Aggregates
- Negation
- Filters
- Binds
- Optionals

SPARQL “used to express queries”



<https://www.w3.org/TR/sparql11-overview/>



W3C Recommendation



SPARQL 1.1 Overview

W3C Recommendation 21 March 2013

This version:

<http://www.w3.org/TR/2013/REC-sparql11-overview-20130321/>

Latest version:

<http://www.w3.org/TR/sparql11-overview/>

Previous version:

<http://www.w3.org/TR/2012/PR-sparql11-overview-20121108/>

Editor:

The W3C SPARQL Working Group, see [Acknowledgements](#) <public-rdf-dawg-comments@w3.org>

Please refer to the [errata](#) for this document, which may include some normative corrections.

See also [translations](#).

Copyright © 2013 W3C® (MIT, ERCIM, Keio, Beihang), All Rights Reserved. W3C [liability](#), [trademark](#) and [document use](#) rules apply.

Query 1 – Basic SELECT



Select queries are used to extract results from the dataset and can be modified to match whatever pattern you desire.

This is the simplest **SELECT** query. It returns all the triples in your dataset.

The **SELECT** keyword determines the type of query.

SELECT simply returns values.

```
1 # Query 1
2
3 # This query matches all the triples,
4 # and returns all the variables from the matched results
5 SELECT ?S ?P ?O
6 WHERE {
7   ?S ?P ?O
8 }
```

The **WHERE** clause describes what the query is looking for.

In this case we want all triples, so we write the variables **?S ?P** and **?O**, showing we want cases where a subject, predicate, and object exist, with any value in any position.

The output variables **?S ?P** and **?O** tell the **SELECT** query what to return from the variables that are found by the **WHERE** clause.

Here we're asking for everything.

We still have to tell the query that the variable **?P** should come from the predicate of a triple, and that we want to consider all triples. Therefore, we still need **?S ?P ?O** inside the **WHERE** clause.



Query 2 – DISTINCT SELECT

Often you'll want to be more specific. This time, we don't want all the triples to be returned, just their predicates.

Using **SELECT** alone would return duplicates, but we don't want that either.

The **DISTINCT** modifier tells the query to disregard duplicate results.

```
1 # Query 2
2
3 # What kinds of edges are in the data?
4 SELECT DISTINCT ?P
5 WHERE {
6   ?S ?P ?O
7 }
```

This time we only ask for **?P** to be returned.

We still have to tell the query that the variable **?P** should come from the predicate of a triple, and that we want to consider all triples. Therefore, we still need **?S ?P ?O** inside the **WHERE** clause.

Query 3 – Finding Classes



We want to SELECT all the DISTINCT classes that exist in our data to find the all the types of things we storing.

In SPARQL, the keyword 'a' is used universally as the predicate for specifying the type of an entity.

```
1 # Query 3
2
3 # What classes of nodes are there in the data?
4 SELECT DISTINCT ?type
5 WHERE {
6   ?S a ?type
7 }
```

We want to find all instances of triples **WHERE** nodes have the 'a' property.

We are only interested in the object of these triples – their class.
We have defined the variable **?type** to help us.

Query 4 – Class Properties



This is one of the most important queries we'll cover today.

We **SELECT** a class of interest, and then return all its properties - not their values, just the properties it possesses.

• The data store holds a list of prefixes. For example, '⋮' is equivalent to <https://rdfox.com/examples/f1/>

```
✓ 9 prefixes in data store

1  # Query 4
2
3  # What edges are there specifically for drivers?
4  SELECT DISTINCT ?P
5  WHERE {
6    ?S a :driver ;
7    ?P ?O .
8  }
```

We want to find all instances of the class `:driver`

Subjects can be shared across multiple lines.

To denote this, all lines except the last must end with a '⋮'.

The last ends with a '⋮'.

• And to consider all triples that share their subjects, `?S`

Exercise 1



➞ Run through **queries 1-4** in the **web console**.

✦ Try them out yourself.
See what they do and how they do it.

Query 5 – Finding Drivers



- ➔ Find all the drivers along with their forename and surname

```
1 # Query 5
2
3 # What are the names of my drivers?
4 SELECT DISTINCT ?driver ?forename ?surname
5 WHERE {
6     ?driver a ___ ;
7     :driver_forename ___ ;
8     :driver_surname ___ .
9 }
```

← We have introduced new variables to the output that must appear in the **WHERE** clause.

Query 6 – A Driver's Properties



- ➔ To build on what we're got so far, let's find all the associated properties and corresponding values of a specific driver. Say, Lewis Hamilton.

```
1 # Query 6
2
3 # Can you return everything about a specific driver, say, Lewis Hamilton?
4 SELECT ?driver ?p ?o
5 WHERE {
6   ?driver a :driver ;
7     --- "Lewis" ;
8     --- "Hamilton" ;
9     --- --- .
10 }
```

Literals, or the values attached to properties, can be specified in queries too.

Here we're looking for a specific string so we use the `"` characters to identify it.

Query 7 – Hamilton's Races



- Find the total number of races Hamilton has taken part in.

```
1  # Query 7
2
3  # Number of races Lewis Hamilton has raced in
4  SELECT (COUNT(?race) AS ?raceCount)
5  WHERE {
6
7      ?driver :driver_forename "Lewis" ;
8              :driver_surname "Hamilton" .
9
10     ?result :result_driver ___ ;
11             :result_race ___ .
12 }
```

The **COUNT** function counts the number of instances of a variable, `?race`, and binds the resulting value as another variable, `?raceCount`.

Query 8 – All Drivers' Races



- ➔ Find the total number of races each driver individually has taken part in.

The **GROUP BY** clause states how the results should be grouped when returned. Aggregate functions, such as **COUNT**, will now be performed for each group.

```
1  # Query 8
2
3  # Number of races per driver, ordered
4  SELECT ?driver (COUNT(?race) AS ?raceCount)
5  WHERE {
6    ?result :result_driver ___ ;
7    :result_race ?race .
8  }
9  GROUP BY ___
10 ORDER BY DESC(___)
```

The **ORDER BY** clause determines the order of the query results.

The **DESC** modifier tells the **ORDER BY** clause the variable indicated should be ranked in descending order.

Exercise 2



➔ Complete **queries 5-8**.



HINTS

It might help you to have **several tabs open** with the previous queries in each one—they will help you fill in the gaps!

You can click on **Incomplete Query** at any time to see an editable version of the query for you to play around with.

Sometimes there will be parts missing for you to fill in. You can click on **Answer** for the query in full for guidance.

Exercise 2



BONUS Qs

Display the **forename and surname** of the driver (**instead of the IRI**) next to their race count.

Answer

Using the **GROUP_CONCAT** aggregate function, find the list of all teammates that each driver has ever had.

Answer



HINTS

Two people are teammates if they competed in the same **race** for the same **constructor**. Notice every result is associated with a constructor and race.

You can put **DISTINCT** inside of an aggregate function to deduplicate.

Query 9 – Drivers Without Podiums



➔ Find all of the drivers who have **never** finished on the podium.

The **FILTER** function contains properties that are used to restrict the results that are returned.

```
1 #Query 9
2
3 # Drivers who never got a podium
4 SELECT --- ---
5 WHERE {
6   # First, get all the drivers with their names.
7     ?driver a :driver ;
8       :driver_forename ?forename ;
9       :driver_surname ?surname .
10
11  # Then make sure that they have never achieved a podium (i.e. positions 3, 2, or 1)
12  → FILTER NOT EXISTS {
13      ?result --- ?driver ;
14      :result_positionOrder ?positionOrder .
15  → FILTER(?positionOrder IN(1, 2, 3))
16  }
17 }
```

The **IN** operator checks to see if the variable on the left-hand side appears in the list.

The **NOT EXISTS** operator is our first glimpse of negation. It returns true if its conditions are **not** matched.

In effect, in this case the conditions of the **FILTER** are reversed, so results are limited to those that **do not** match what is stated.

Query 10 – Most Races Without a Podium



- ➔ Rank all of the drivers who have **never** finished on the podium from most to least races without a podium.

Here we have used a query within a query—an ‘inner query’, or ‘subquery’.

```
1 # Query 10
2
3 # Most races run without ever getting a podium
4 SELECT ?forename ?surname ?raceCount
5 WHERE {
6   # First, get all the drivers with their names.
7   # The order of the query atoms is mostly not important, as RDFS will automatically
8   # rearrange it internally in order to have the best performance.
9   ?driver a :driver ;
10      :driver_forename ?forename ;
11      :driver_surname ?surname .
12
13   # Then make sure that they have never achieved a podium (i.e. positions 3, 2, or 1)
14   FILTER NOT EXISTS {
15     ?result2 :result_driver ?driver ;
16     :result_positionOrder ?positionOrder .
17     FILTER(?positionOrder IN (1, 2, 3))
18   }
19
20   # Finally find out how many races they have raced in
21   # This is done in an 'inner query'
22   {
23     SELECT ___ (COUNT(___)) AS ?raceCount
24     WHERE {
25       ?result :result_driver ?driver ;
26       :result_race ?race .
27     }
28     GROUP BY ___
29   }
30 }
31 ORDER BY DESC(___)
```

Query 11 – Driver Win Percentage



- ➔ Calculate the win percentage of each driver individually, ranking them from most to least successful.

The **BIND** function sets a variable to a specific value.

```
1  # Query 11
2
3  # Drivers with their win percentage, ordered by win percentage
4  SELECT ?forename ?surname ?raceCount ?raceWins ?percentage
5  WHERE {
6    # First get the drivers
7    ?driver :driver_forename ?forename ;
8           :driver_surname ?surname .
9
10 # Then get the race count for each driver with an innery query...
11 {
12     SELECT ?driver (COUNT(?race) AS ?raceCount)
13     WHERE {
14         ?result :result_driver ?driver ;
15                :result_race ?race .
16     }
17     GROUP BY ?driver
18 }
19
20 # ... and get the *win* count for each driver with another inner query.
21 {
22     SELECT ___ (COUNT(?race) AS ___)
23     WHERE {
24         ?result :result_driver ___ ;
25                :result_race ___ ;
26                :result_positionOrder 1 .
27     }
28     GROUP BY ?driver
29 }
30
31 # Finally use the two aggregate variables to compute a percentage
32 # with the BIND keyword.
33 BIND(?raceWins/___ AS ?percentage)
34 }
35 ORDER BY DESC(?percentage)
```

Query 12 – Correcting Win Percentage



- Correct the calculation of win percentages to include even those drivers who never won.

The **OPTIONAL** keyword states that if, for a given case, the subsequent result cannot be evaluated, skip over it and leave it undefined.

The **COALESCE** function returns the first value in its list that does not throw up an error (including undefined).

It is bound to the stated variable.

```
1 # Query 12
2
3 # Can you spot a problem?
4
5 # We only return 'biased' results, i.e. those who have at least one win!
6 # We will solve this using the so called OPTIONAL keyword
7
8 SELECT ?forename ?surname ?raceCount ?raceWinsFinal ?percentage
9 WHERE {
10 # First get the drivers
11 ?driver :driver_forename ?forename ;
12         :driver_surname ?surname .
13
14 # First get the race count for each driver.
15 {SELECT ?driver (COUNT(?race) AS ?raceCount)
16  WHERE {
17    ?result :result_driver ?driver ;
18            :result_race ?race .
19  }
20  GROUP BY ?driver}
21
22 # Then *optionally* get the win count for each driver.
23 OPTIONAL {SELECT ?driver (COUNT(?race) AS ?raceWins)
24  WHERE {
25    ?result :result_driver ?driver ;
26            :result_race ?race ;
27            :result_positionOrder 1 .
28  }
29  GROUP BY ?driver}
30
31 # Using the COALESCE keyword, we make sure that when the number of wins is undefined
32 # it is 'bound' as 0.
33 BIND(COALESCE(____,0) AS ?raceWinsFinal)
34 # And use the two aggregate variables to compute a percentage
35 BIND(____ / ?raceCount AS ?percentage)
36
37 }
38 ORDER BY DESC(____)
```

Exercise 3



➔ Complete **queries 9-12**.

i HINTS

To finish on the podium you must come in positions 1, 2, or 3.

Every result will be associated with a position number which tells us where they finished in the race.

Exercise 3 – Bonus Question



- Use **FILTER** to restrict this list to only drivers who **won at least 5** races.
- Answer

Running SPARQL from the command line



<https://docs.oxfordsemantic.tech/command-line-reference.html>

As well as running SPARQL queries from the web console, queries can be run from the command line.

Here are some of the commands that you will need to set this up:

```
active <data store name>
```

The active command sets the active data store for subsequent commands.

```
set output out
```

The **set** command allows you to change internal variables used by the RDFox shell.

set output out sets the output variable so that query results are sent to the terminal.

```
evaluate <query file name>
```

Runs the query with that filename.

Running SPARQL from the command line



From the RDFox command line:

- 1 Make the **f1** data store active. `active f1`
- 2 Set the output to the terminal. `set output out`
- 3 Then specify a prefix. `prefix : <https://rdfox.com/examples/f1/>`
- 4 Run **query 6** again.
If you are running RDFox from `RDFoxWorkshop` directory, then the command should be
`evaluate queries/q6.rq`

Congratulations!



A look back at what you've achieved.

The Foundations

- ✓ What is a database?
- ✓ What is a triple?
- ✓ What is a Knowledge Graph?
- ✓ What is Semantic Reasoning?
- ✓ What is RDFox?

Setting up RDFox

- ✓ License and executable
- ✓ VS Code
- ✓ REST endpoint and UI
- ✓ Create a data store
- ✓ Import data

Querying with SPARQL

- ✓ Basics of SPARQL
- ✓ Useful queries to explore
- ✓ Aggregates
- ✓ Negation
- ✓ Filters
- ✓ Binds
- ✓ Optionals

Closing RDFox



If you are continuing with the RDFox Advanced Reasoning Workshop, please keep RDFox running.

Otherwise, stop your RDFox server with the `quit` command.

Additional information



Visit our website

<https://www.oxfordsemantic.tech>



Our blog

[https://www.oxfordsemantic.tech/
blog](https://www.oxfordsemantic.tech/blog)



Read our documentation

[https://docs.oxfordsemantic.tech/
index.html](https://docs.oxfordsemantic.tech/index.html)



Request an evaluation license

[https://www.oxfordsemantic.tech/
free-trial](https://www.oxfordsemantic.tech/free-trial)